

Calibre / Cepheid

Reviewable fixes and the cost of session triage

A refreshed fact-check companion to edition six, linked to edition seven's discussion of verification value. This is a concise review brief, not a full article reproduction.

The account we want to preserve

Cepheid was presented as Calibre's internal maintenance agent. Session evidence feeds a cheap triage pass; findings can escalate to a coding agent. A proposed fix carries three videos: the original session, reproduction and verification. More ambitious changes retain an engineer in the loop. Please correct that scope if it is wrong.

Cost figures and review questions

The event materials give \$11.94 per 1,000 replays for the Luna triage case. The \$4.82 GLM comparison assumes the same usage and includes promotional pricing; it should not read as a measured replacement run or a current quote. Stronger-model fixes and hosting are additional.

Please confirm:

1. Which cost rows were measured, estimated or calculated from a rate card?
2. We read 38.298M uncached input, 97.417M cached input and 1.942M output tokens. Are those counts and the input-cache denominator correct?
3. What proportion of sessions escalates, and what is the resulting full cost including fixes, hosting and review?
4. Are the internal-agent description, photographs and three-video workflow accurate?

Articles and response

Edition six: A computer your agent can throw away

Edition seven: The gate had no buyer

Please identify the sentence or number you would correct and the supporting source. Clean slide exports are welcome. Any proposed quotation will be checked with you before publication. Antonio Serrano / pulse@autopilotindex.com

Cost provenance

Cepheid: cost with Luna				Calibre
Triage cost for 1,000 replays \$11.94				
Token type	Tokens	Price / 1M	Cost	
Uncached input	38.298M	\$0.20	\$7.66	
Cached input	97.417M	\$0.02	\$1.95	
Output	1.942M	\$1.20	\$2.33	
Estimate from recorded triage usage. Stronger-model fixes and hosting are additional.				

Original cost-slide photograph from the edition-six reporting materials. Rates and workload assumptions are historical, not a current pricing recommendation.

Derived scenario - reproduced from the slide

Total input is 135.715M tokens. Keeping output at 1.942M tokens and using the photographed rates:
 $135.715 \times 10\% \times \$0.20 + 135.715 \times 90\% \times \$0.02 + 1.942 \times \$1.20 = \7.48757 , or **\$7.49.**

This reproduces the published estimate. The starting input-cache rate is about 71.78%, rounded to 72%. It is a scenario calculation, not an observed improvement.

The cache scenario saves about 37% against \$11.94. The quoted \$4.82 model-price scenario saves about 60%. Both are material, but the latter is not an observed quality-equivalent replacement run. Neither is the full cost of operating Cepheid.

Link to edition seven

The three-video pull request is an example of verification producing evidence a reviewer can use. That connection is the author's interpretation, not a claim by Calibre about the broader evaluation market.