

Veris AI

Fidelity evidence, coverage and the testing-to-training question

A refreshed fact-check companion to editions six and seven. This is a concise review brief, not a full reproduction of the published sections. The existing event photographs and article remain the reporting record.

Published evidence - correction

The fidelity figures are published at veris.ai/fidelity. Veris reports 703/721 exact matches across 15 services in Rev 2.0, with 18 mismatches, four attributed to the twins, and no false passes in that measured population. These are vendor-run results using an external test suite, not our independent replication. They do not establish fidelity for the entire catalogue.

What we would like you to check

1. Is Rev 2.0 the appropriate report to cite? Can readers obtain a pinned test-suite revision and the operation-level evidence needed to reproduce it?
2. How do customers identify unsupported operations, and how are vendor API changes reflected in twin versions?
3. The event account describes the same sanctions-screening task finishing in 32 minutes with missed matches, versus 40 minutes with all four caught. Which agent/model versions, prompts and repetitions underpin that comparison?
4. Edition seven describes a move from testing toward training. Is that an accurate change of emphasis, or should it describe complementary uses?


Articles and response

Edition six: [A computer your agent can throw away](#)

Edition seven: [The gate had no buyer](#)

Please identify the sentence or number you would correct and the supporting source. Clean slide exports are welcome. Any proposed quotation will be checked with you before publication. Antonio Serrano / pulse@autopilotindex.com

Demonstration context



Photograph from the edition-six reporting materials. The demo is evidence about this task and setup, not a general performance guarantee.

Proposed wording for edition seven

"The same environment can support verification and training. My interpretation is that its value grows when it helps an agent learn or complete work, rather than merely blocking a release. Whether that describes Veris's commercial direction is a question for the company."

Review boundary

The reporting should separate what the demonstration showed, what the benchmark reports and what the author infers about the business. Confirmation of factual details would not imply endorsement of the article's market thesis.

The original review pack lacked the direct public-report link and clear replication scope. This version supplies both. It does not describe the figures as unpublished.